

COCO Dataset - use Faster RCNN + MobileNet to conduct Object Detection

Created 2023-11-03 | Updated 2023-12-10 | Code > Machine Learning

| Word Count: 3.1k | Reading Time: 19mins | Post Views: 0 | Comments: 0

💡 Introduction

Recently I took an AI course, the main content is the following topics:

1. Learn about Coco dataset
2. User pre-trained version of Faster R-CNN to predict the bounding box
3. Calculate IoU

💡 Homework Requirement

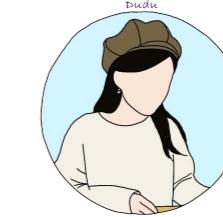
1. **Download the Coco Collection** *: download the files "2017 Val images [5/1GB]" and "2017 Train/Val annotations [241MB]" from the Coco page.

Download from Coco page. You can load them into your notebook using the pycocotools library.

2. **Randomly select ten from the dataset** : 10 images are randomly selected from this dataset.

3. **Predict the box using the pre-trained model FasterR-CNN** : use a pre-trained version of the Faster R-CNN (Resnet50 backbone) to predict the bounding box of the object on the 10 images. Only regions with scores greater than 0.8 are retained.

4. **Visualize the model together with the answer visualization** : visualize the predicted bounding boxes



Shannon Hung

In order to avoid forgetfulness, I started to record

Articles	Tags	Categories
13	12	12

Follow Me



📢 Announcement

I am a graduate student at the Taiwan University of Science and Technology - Institute of Information Management, planning to pursue a dual-degree program in Germany this year. The

boxes and label. Show all 10 pairs of images side by side in the jupyter notebook.

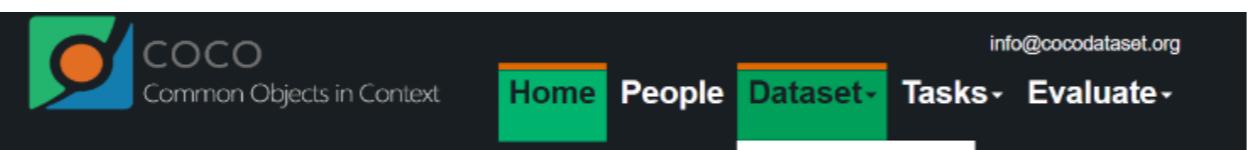
5. **Use another pre-trained model Mobilnet** : Repeat the above steps using the Mobilenet backbone of the Faster R-CNN.
6. **Calculate IoU Compare Models** : Which backbone provides better results? Calculate the IoU for both methods.

Task 1: Downloading the COCO Dataset

Task 1

- i 1. **Download the COCO Dataset** : Obtain the files "2017 Val images [5/1GB]" and "2017 Train/Val annotations [241MB]" from the Coco page. Utilize the pycocotools library to import them into your notebook.

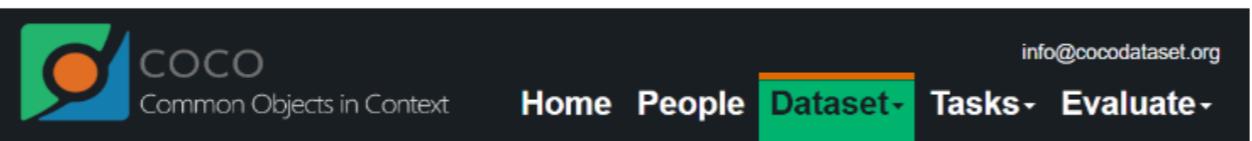
You can follow this guide to proceed with the download: Download COCO Dataset



News

- We are pleased to announce the COCO 2020 DensePose Challenges.
- The new rules and awards for this year challenges encourage innovative methods.
- Results to be announced at the [Joint COCO and LVIS Recognition ECCV workshop](#).

這部分一樣，就連到他們官網後，點 Dataset 然後選 Download，接著會進入下載頁面：



Tools

COCO API

Images

- 2014 Train images [83K/13GB]
- 2014 Val images [41K/6GB]
- 2014 Test images [41K/6GB]
- 2015 Test Images [81K/12GB]
- 2017 Train Images [118K/18GB] (highlighted)
- 2017 Val Images [5K/1GB]
- 2017 Test Images [41K/6GB]
- 2017 Unlabeled Images [123K/19GB]

Annotations

- 2014 Train/Val annotations [241MB]
- 2014 Testing Image info [1MB]
- 2015 Testing Image info [2MB]
- 2017 Train/Val annotations [241MB] (highlighted)
- 2017 Stuff Train/Val annotations [1.1GB]
- 2017 Panoptic Train/Val annotations [821MB]
- 2017 Testing Image info [1MB]
- 2017 Unlabeled Image info [4MB]

BASH

```
1 .
2 |   annotations # These are annotation files
```

have learned a lot, I tend to forget much of it. To combat forgetfulness, I have embarked on my note-taking journey.

Contents

24

Introduction

Homework Requirement

Task 1: Downloading the COCO Dataset

Task 2: Randomly Select Ten Images

Setting up the COCO API

Annotation Visualization

Randomly Select Ten Images

Task 3+5: FasterR-CNN v.s Mobilnet

Using pre-train model

Convert image to tensor

Training the model

Only select the prediction results > 0.8

Task 4+6: Visualization + IoU

Supplement: IoU

Recent Post



Windows - A Step-by-Step Guide to Enable CUDA GPU on Ten...
2024-09-04



How to Handle Multi-Column Text Sorting with Amazon Textract
2024-03-20



Twitter Dataset - Using LSTM to predict the emotion of the arti...
2023-11-16



COCO Dataset - use Faster RCN N + MobileNet to conduct Ob...

```
6   |   └── instances_val2017.json
7   |   └── person_keypoints_train2017.json
8   |   └── person_keypoints_val2017.json
9   └── val2017 # This is the image set
10  |   ├── 00000000139.jpg
11  |   ├── 00000000285.jpg
12  |   ├── 00000000632.jpg
13  |   ├── 00000000724.jpg
14  |   ├── 00000000776.jpg
15  |   ├── 00000000785.jpg
16  |   ├── 00000000802.jpg
17  |   ...
```



Flower102 Dataset - Using Transfer Learning to train + Using ...
2023-10-31

- Download these two files as shown in the image.
- After downloading, the folder structure upon extraction will resemble the one above.

⚡ Task 2: Randomly Select Ten Images

Task 2

2. **Randomly Select Ten Images from the Dataset** : Pick 10 images randomly from this dataset.

Here, we'll primarily do a few things:

- Import necessary libraries.
- Set up the COCO API to allow it to access relevant information from our dataset, such as bounding box positions, label locations, and image information.
- Visualize images and perform annotations.
- Randomly select ten images.

Let's begin by importing the necessary libraries.

```
● ● PYTHON
1
2 # CNN
3 import torch.nn.functional as F
4 import torch
5 import torch.nn as nn
6 import torch.optim as optim
7 from torch.optim import lr_scheduler
8 import torch.backends.cudnn as cudnn
```

```
12 import matplotlib.pyplot as plt
13 import time
14 import os
15 from PIL import Image
16 from tempfile import TemporaryDirectory
17 import time
18 import random
19
20 # torchvision
21 import torchvision
22 import torchvision.transforms as transforms
23
24 # dataset
25 from pycocotools.coco import COCO
26 import cv2
27
28
29 cudnn.benchmark = True
30 plt.ion() # interactive mode
```

⚡ Setting up the COCO API

COCO provides an API to access datasets. By providing it with a JSON file, we can easily retrieve the necessary information such as images, labels, bounding boxes, and more.

PYTHON

```
1 # Specify dataset location
2 cocoRoot = "../../Data/Coco/"
3 dataType = "val2017"
4
5 # Set annotation file location
6 annFile = os.path.join(cocoRoot, f'annotations/instances_{dataType}.json')
7 print(f'Annotation file: {annFile}')
8 # # initialize COCO api for instance annotations
9 coco=COCO(annFile)
10 coco
```

Result

LUA

```
1 Annotation file: ../../Data/Coco/annotations/instances_val2017.json
2 -- Indicates successful annotation file read
3 loading annotations into memory...
```

```
5 creating index...
6 index created!
```

.Annotation Visualization

To ensure familiarity with the COCO-provided API, here's an exercise focusing on the following:

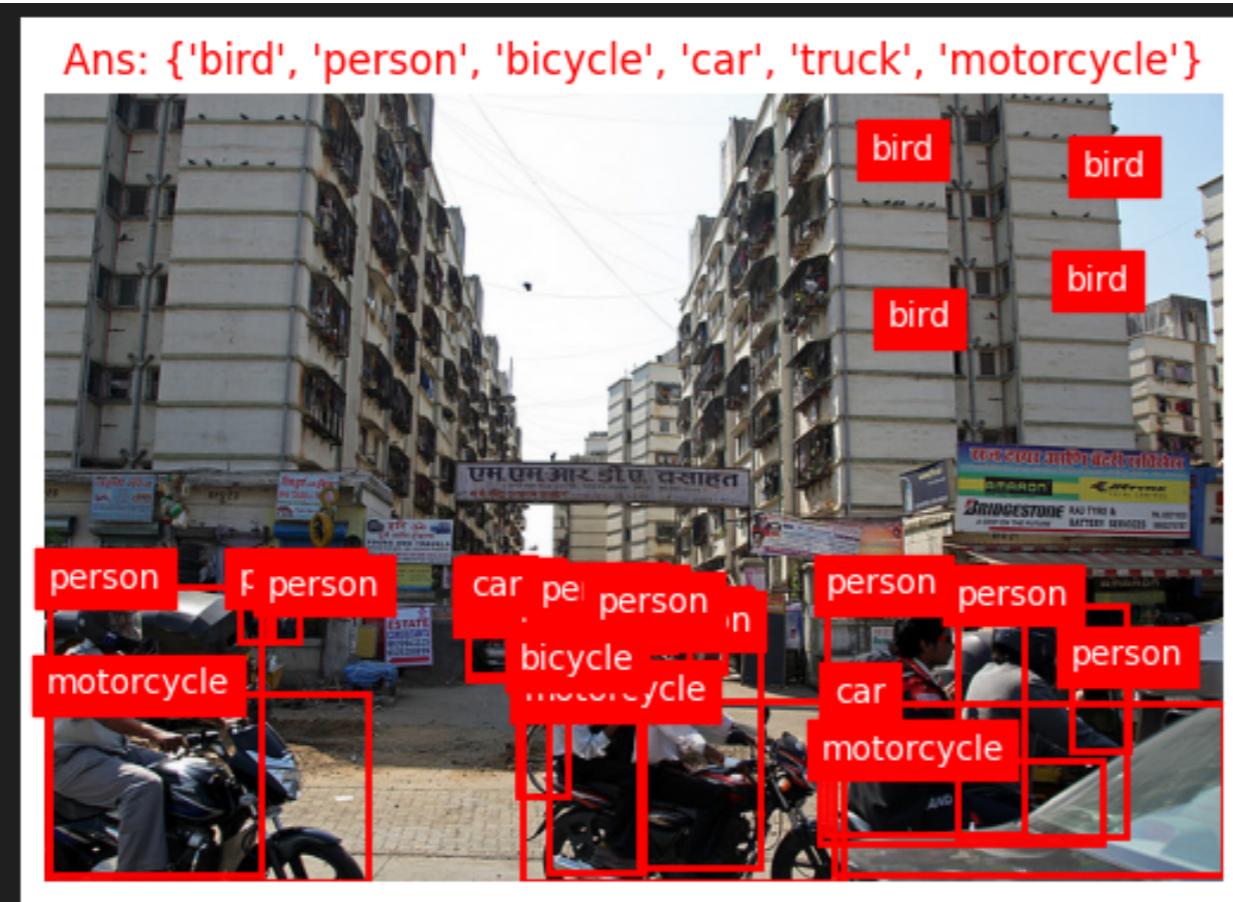
- Obtaining image info by ID
- Retrieving annotation info by ID
- Learning to draw bounding boxes and labels on images

PYTHON

```
1 import matplotlib.pyplot as plt
2 from matplotlib.patches import Rectangle
3
4 # Create a function that, given an image ID, plots the image with bounding boxes and labels
5 def plot_image_with_annotations(coco, cocoRoot, dataType, imgId, ax=None):
6     # Get image information
7     imgInfo = coco.loadImgs(imgId)[0]
8     # Get image location for visualization
9     imPath = os.path.join(cocoRoot, dataType, imgInfo['file_name'])
10    # Read the image
11    im = cv2.imread(imPath)
12    # Convert color space: OpenCV defaults to BGR, but matplotlib to RGB, so conversion is needed
13    im = cv2.cvtColor(im, cv2.COLOR_BGR2RGB)
14
15    # Find all annotations (bounding boxes) for the image
16    annIds = coco.getAnnIds(imgIds=imgInfo['id'])
17    # Load all annotation information: bounding box coordinates, labels, accuracies
18    anns = coco.loadAnns(annIds)
19    all_labels = set()
20
21    # Extract bounding box coordinates, labels, accuracies
22    for ann in anns:
23        # Specifically select information related to the bounding box: returns (x, y) of the top-left corner, width, height
24        x, y, w, h = ann['bbox']
25
26        # Get label text information: load category name by category ID
27        label = coco.loadCats(ann['category_id'])[0]["name"]
28        all_labels.add(label)
29
30        # Draw bounding boxes using provided coordinates
31        rect = Rectangle((x, y), w, h, linewidth=2, edgecolor='r', facecolor='none')
32
33    # Draw the image: if sorting of images is needed, ax parameter specifies the position
```

```
36         plt.text(x, y, f'{label}', fontsize=10, color='w', backgroundcolor='r')
37     else:
38         ax.add_patch(rect)
39         ax.text(x, y, f'{label}', fontsize=10, color='w', backgroundcolor='r')
40
41     # Display the image with a title
42     if ax is None:
43         plt.imshow(im)
44         plt.axis('off')
45         plt.title(f'Annotations: {all_labels}', color='r')
46         plt.show()
47     else:
48         ax.axis('off')
49         ax.set_title(f'Annotations: {all_labels}', color='r', loc='center', pad=20)
50         ax.imshow(im)
51
52
53 # Get the tenth image
54 imgIds = coco.getImgIds()
55 imgId = imgIds[10]
56 # Plot the image with annotations
57 plot_image_with_annotations(coco, cocoRoot, dataType, imgId)
```

Result

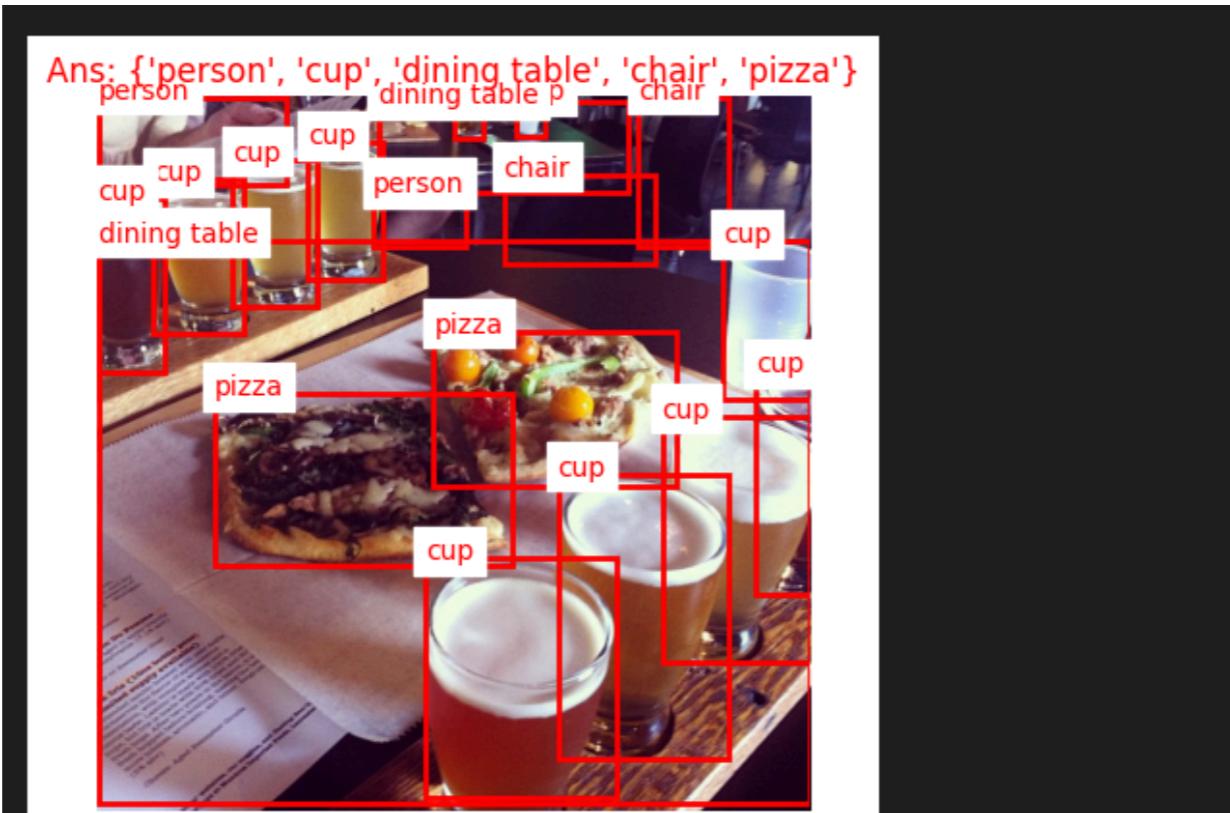


⚡ Randomly Select Ten Images

🟡 ● PYTHON

```
1 def random_select(coco, cocoRoot, dataType, num_images=10):
2     # Get IDs for all images
3     imgIds = coco.getImgIds()
4     # Randomly select num_images IDs from this set
5     selected_imgIds = random.sample(imgIds, num_images)
6     # Call the plot_image_with_annotations function for each selected ID
7     for imgId in selected_imgIds:
8         # Plot images based on their IDs
9         plot_image_with_annotations(coco, cocoRoot, dataType, imgId)
10
11    # Print out all selected IDs
12    return selected_imgIds
13
14 valid_ids = random_select(coco, cocoRoot, dataType, num_images=10)
15 valid_ids
```

Result



[385719, 210388, 481413, 557258, 328238, 531134, 300155, 18193, 218362, 74733]

⚡ Task 2 : E. FasterR-CNN vs Mobilenet

Task 3 & 5

3. **Predicting bboxes using the pre-trained model FasterR-CNN** : Use a pre-trained version of Faster R-CNN (Resnet50 backbone) to predict the bounding box of objects on the 10 images. Only keep regions that have a score > 0.8.
5. **Using another pre-trained model Mobilnet** : Repeat the steps from above using a Mobilenet backbone for the Faster R-CNN.

⚡ Using pre-train model

🟡 ● PYTHON

```
1 # using pre-train model (FasterR-CNN)
2 model_res = torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="FasterRCNN_ResNet50_
3 model_res.eval()
4
5 # using pre-train model (Mobilenet)
6 model_mobile = torchvision.models.detection.fasterrcnn_mobilenet_v3_large_fpn(weights=torchvisi
7 model_mobile.eval()
```

⚡ Convert image to tensor

We need to be able to take the image out of the picture based on the position of the image. Then we need to convert the image read from the book into a tensor before we can put it into the model for prediction. So we made two functions:

- One is to read the picture
- One is to convert the picture to a tensor.

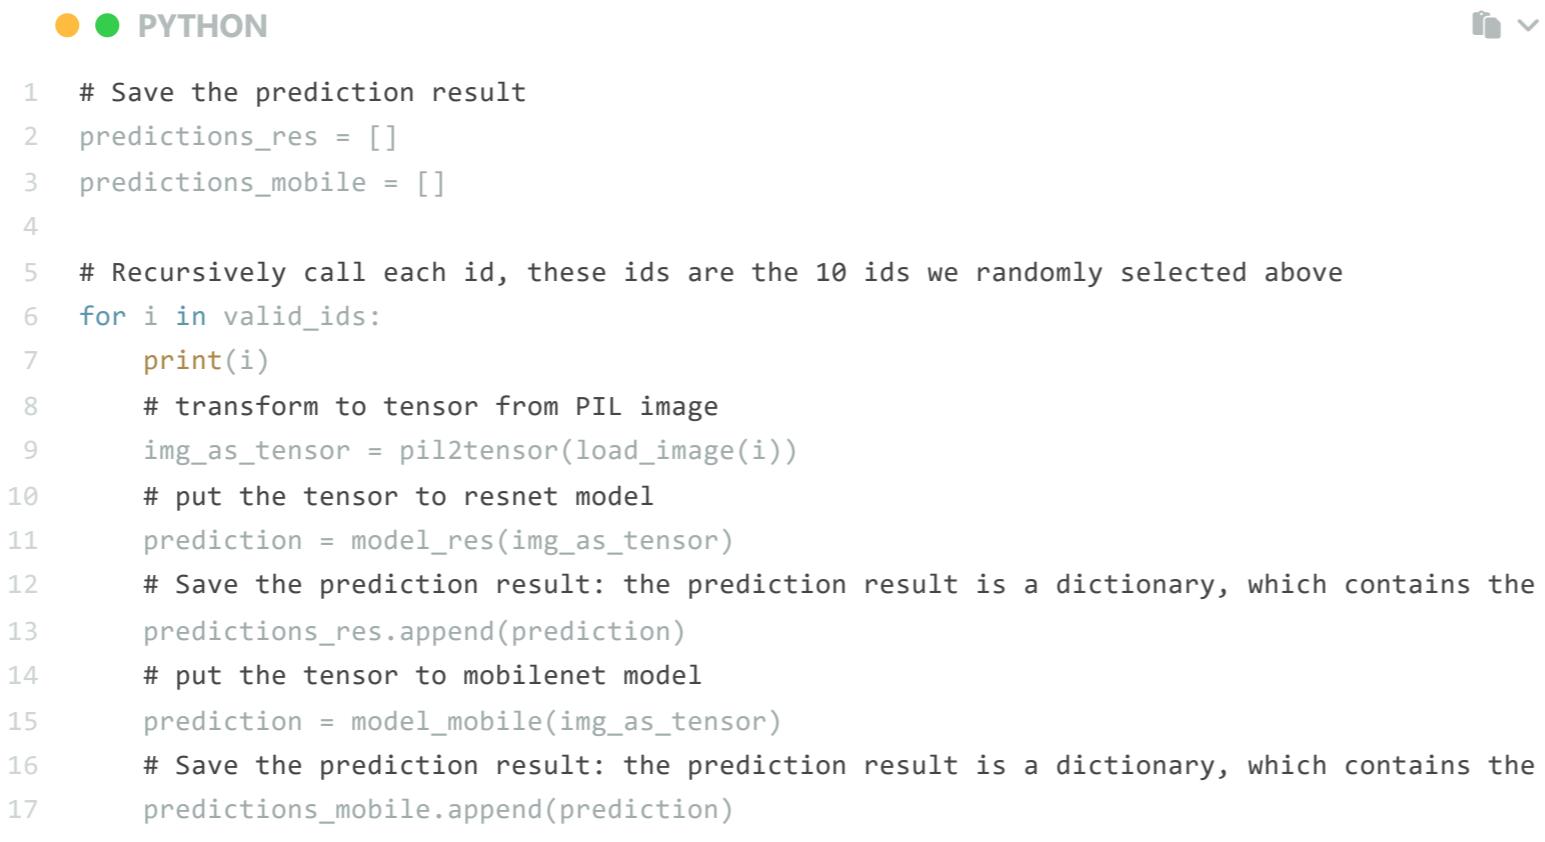
🟡 ● PYTHON

```
1 from PIL import Image
2
3 def load_image(imgIdx):
4     # Get image information
5     imgInfo = coco.loadImgs(imgIdx)[0]
6     # Get image location for visualization
7     imPath = os.path.join(cocoRoot, dataType, imgInfo['file_name'])
8     # Read the image path
9     print(imPath)
10    try:
11        # Read the image
```

```
14     raise Exception()
15
16 # Convert the picture to a tensor
17 def pil2tensor(pil_image):
18     # Use unsqueeze(0) because the model still contains the batch size dimension, a total of
19     # But the picture has only one picture without batch size, the picture is converted to te
20     # /255 is because the input of the model is a number between 0 and 1, and the value of th
21     return torchvision.transforms.PILToTensor()(pil_image).unsqueeze(0) / 255.0
```

👉 Training the model

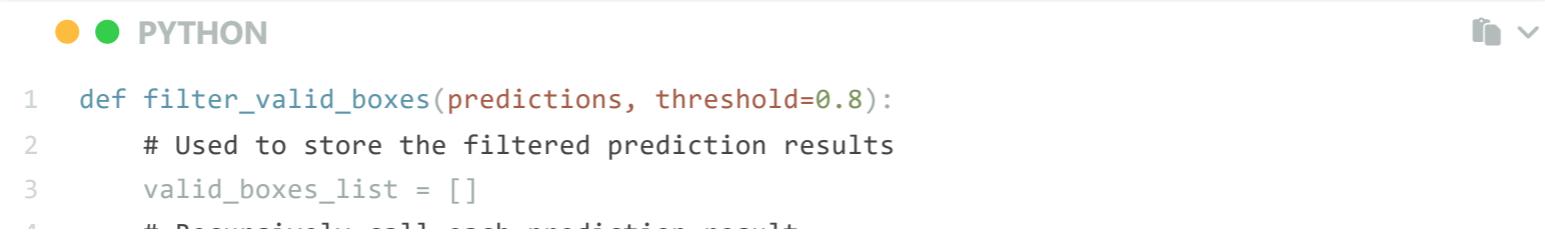
After the pre-training model is loaded, we need to train the model. The training process is as follows:



```
● ● PYTHON
1 # Save the prediction result
2 predictions_res = []
3 predictions_mobile = []
4
5 # Recursively call each id, these ids are the 10 ids we randomly selected above
6 for i in valid_ids:
7     print(i)
8     # transform to tensor from PIL image
9     img_as_tensor = pil2tensor(load_image(i))
10    # put the tensor to resnet model
11    prediction = model_res(img_as_tensor)
12    # Save the prediction result: the prediction result is a dictionary, which contains the p
13    predictions_res.append(prediction)
14    # put the tensor to mobilenet model
15    prediction = model_mobile(img_as_tensor)
16    # Save the prediction result: the prediction result is a dictionary, which contains the p
17    predictions_mobile.append(prediction)
```

👉 Only select the prediction results > 0.8

After the model is trained, we need to select the prediction results that are greater than 0.8. The reason is that the model will predict a lot of bounding boxes, but we only need the bounding boxes with high accuracy. So we need to filter out the bounding boxes with low accuracy. The code is as follows:



```
● ● PYTHON
1 def filter_valid_boxes(predictions, threshold=0.8):
2     # Used to store the filtered prediction results
3     valid_boxes_list = []
4     # Remove all the prediction results whose confidence score is less than the threshold
```

```

7     # Recursively call each bounding box
8     for box, label, score in zip(prediction[0]["boxes"], prediction[0]["labels"], prediction[0]["scores"]):
9         # Only keep the predicted bounding box with accuracy greater than threshold
10        if score >= threshold:
11            # Save the predicted bounding box, label, and accuracy
12            valid_boxes_for_this_prediction.append((box, label, score))
13        # If none of the predicted boxes in this picture have an accuracy greater than threshold
14        valid_boxes_list.append(valid_boxes_for_this_prediction)
15    # Return the filtered prediction result
16    return valid_boxes_list
17
18 # Set threshold to 0.8 and get the prediction results of resnet and mobilenet
19 valid_boxes_res = filter_valid_boxes(predictions_res, threshold=0.8)
20 valid_boxes_mobile = filter_valid_boxes(predictions_mobile, threshold=0.8)

```

📌 Task 4+6: Visualization + IoU

Tasks 4 & 6

1. **Visualize the model together with the solution** : Visualize the predicted bounding boxes and label together with the ground truth bounding
2. **CalculateIoU to compare models** : Which backbone delivers the better results? Calculate the IoU for both approaches.

There are a few important points in visual dialog, the steps are as follows:

- We need to know the id of the image first, and get the annotation information based on the id, then we can Calculate the IoU.
- We take the annotation information and the model information to conduct the IoU Calculate.
- We read the location of the image in the computer, and according to the path of the image, we draw the image through plt first.
- We read the location of the picture in the computer, and based on the path of the picture, we draw the picture through plt first, and then based on the picture, we can draw the prediction box and label on it, as well as the average value of the IoU.

The following program is the procedure described above, we will draw the results of both models and calculate the average of the IoU.

 PYTHON



```

1 import matplotlib.pyplot as plt
2 from PIL import Image
3 import os
4
5
6 # Can put the results of different models into this function, and it will return the average
7 def display_annotated_results(imgId, valid_boxes, model_name, color='g', ax=None):
8     # Load the image
9     imgInfo = coco.loadImgs(imgId)[0]
10    image_path = os.path.join(cocoRoot, dataType, imgInfo['file_name'])
11    image = Image.open(image_path)
12
13    # Get the correct bounding box results
14    annIds = coco.getAnnIds(imgIds=imgInfo['id'])
15    anns = coco.loadAnns(annIds)
16    bbox_tlist_anns = torch.tensor([ann["bbox"] for ann in anns]) # tensor.shape[2,4]
17    # because our bounding box is x,y,w,h which is the coordinate of the lower left corner of
18    # But torchvision Calculate the box_iou must give the coordinates of the lower left corner
19    # x,y,w,h -> x1,y1,x2,y2 = x,y,x+w,y+h
20    bbox_tlist_anns[:, 2] = bbox_tlist_anns[:, 0] + bbox_tlist_anns[:, 2]
21    bbox_tlist_anns[:, 3] = bbox_tlist_anns[:, 1] + bbox_tlist_anns[:, 3]
22
23    # From resultsvalid_boxes, we only need the box part, so we use (box, _, _)
24    # Use stack because we want to stack all the boxes together to become a tensor
25    bbox_tlist_model = torch.stack([box for box, _, _ in valid_boxes]) # turn [4] to tensor.s
26    # use box_iou 来Calculate IoU
27    iou = torchvision.ops.box_iou(bbox_tlist_anns, bbox_tlist_model) # get IoU
28    # Get the maximum value of each predicted box in ann, and then Calculate the average value
29    avg_iou = np.mean([t.cpu().detach().numpy().max() for t in iou]) # calculate the mean of
30
31    # display image label
32    all_labels = set()
33
34    # Start drawing the prediction box
35    for boxes in valid_boxes:
36        # Get the information of the prediction box, including box, label, and accuracy
37        box, label, score = boxes
38        # Get the text information of the label: load category name by category ID
39        label = coco.loadCats(label.item())[0]["name"]
40        # Save the label for later display
41        all_labels.add(label)
42        # Because the results returned by the model are two coordinates, the lower left corner
43        x, y, x2, y2 = box.detach().numpy() # x,y,w,h -> x,y,x2-x,y2-y
44        rect = Rectangle((x, y), x2 - x, y2 - y, linewidth=2, edgecolor=color, facecolor='none')
45
46        # Draw the picture: if you need to sort the picture, you can specify where to draw the
47        if ax is None:

```

```

50     # Draw the label on the prediction box
51     plt.text(x, y, f'{label}', fontsize=10, color='w', backgroundcolor=color)
52 else:
53     # add_patch can add a patch to the current axes, and then draw the prediction box
54     ax.add_patch(rect)
55     # Draw the label on the prediction box
56     ax.text(x, y, f'{label}', fontsize=10, color='w', backgroundcolor=color)
57
58 # display image and give it a title, the title is the label that appears in this picture,
59 if ax is None:
60     plt.axis('off')
61     plt.title(f'{model_name}: {all_labels} \n IoU: {avg_iou:.4f}', color=color)
62     plt.imshow(image)
63     plt.show()
64
65 else:
66     ax.axis('off')
67     ax.set_title(f'{model_name}: {all_labels} \n IoU: {avg_iou:.4f}', color=color)
68     ax.imshow(image)
69
70 return avg_iou
71
72
73 res_iou = []
74 mobile_iou = []
75
76 # Recursively call each id, where id is one of the 10 random ids we selected above
77 for i in range(len(valid_ids)):
78     # Create a 1x3 grid of images, each image sized 15x5
79     fig, axs = plt.subplots(1, 3, figsize=(15, 5))
80
81     # Draw the truth image and display it in the center
82     plot_image_with_annotations(coco, cocoRoot, dataType, valid_ids[i], ax=axs[1])
83
84     # Draw the predicted results from two different models on the left and right sides respectively
85     i_mobil_iou = display_annotated_results(valid_ids[i], valid_boxes_mobile[i], "mobile", color='red')
86     i_res_iou = display_annotated_results(valid_ids[i], valid_boxes_res[i], "ResNet", color='blue')
87
88     # Store the IoU of each image to assess the overall performance of the model
89     mobile_iou.append(i_mobil_iou)
90     res_iou.append(i_res_iou)
91
92     # Organize the layout
93     plt.tight_layout()
94
95 # Print the mean of the IoU list
96 print("ResNet: Avg.", np.mean(res_iou), "; each IoU:", res_iou)
97 print("MobileNet: Avg." , np.mean(mobile_iou), "; each IoU:" , mobile_iou)

```

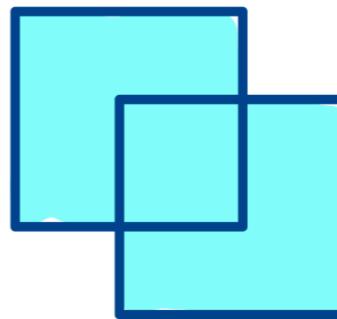
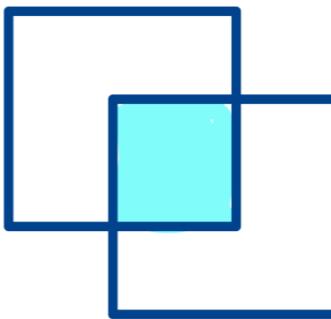


⚡ Supplement: IoU

- Ref: <https://blog.csdn.net/lAMoldpan/article/details/78799857>
- Ref: <https://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/>

i IoU (Intersection over Union) is a metric used to evaluate object detection algorithms. It is defined as the `intersection area` of the `predicted box` and the `true box` divided by their `union area`. The value ranges between 0 and 1, where a higher value indicates a greater overlap between the predicted and true boxes, signifying more accurate predictions.

$$IoU = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$



Source: <https://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/>

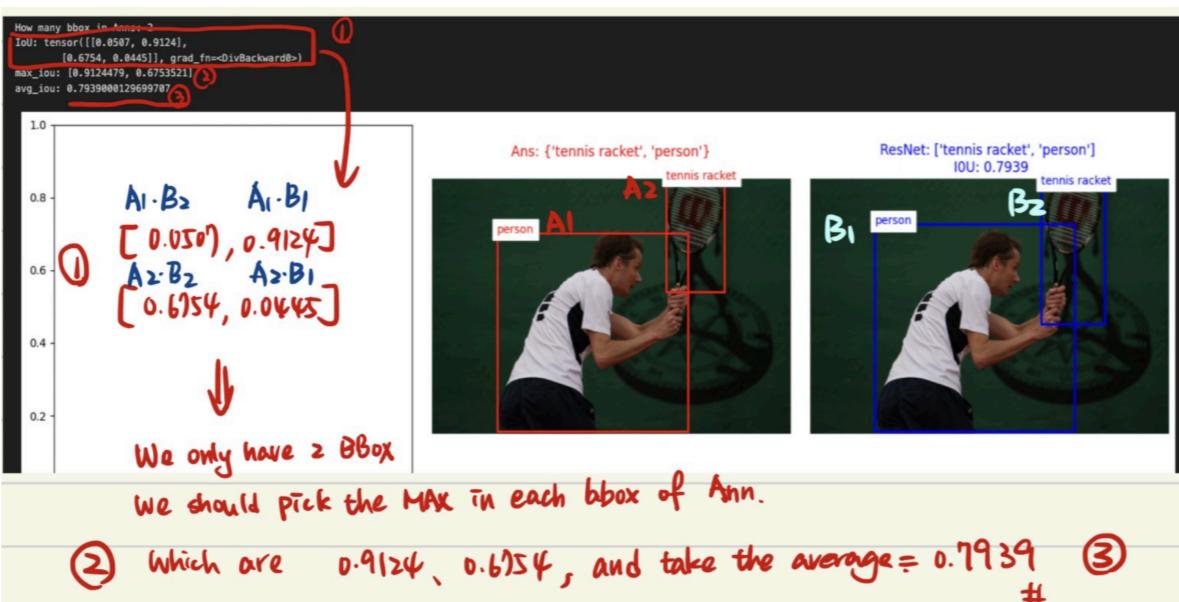
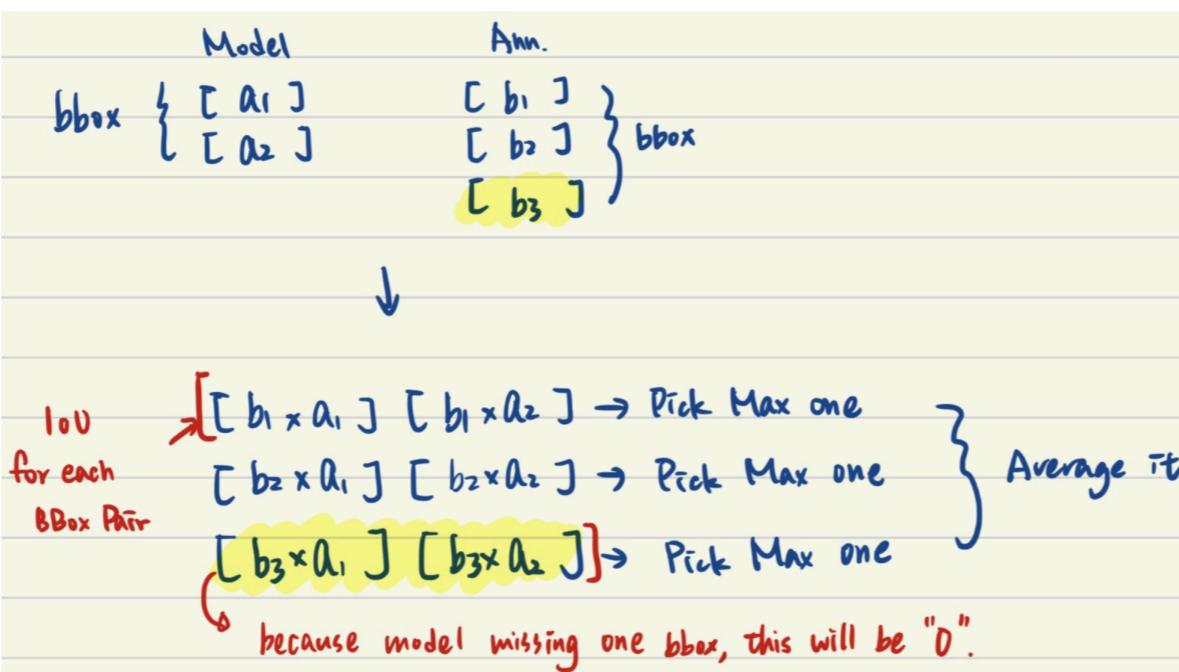
i From the example above, you might be wondering, what exactly does the following code segment do?

● ● PYTHON

```
1 torchvision.ops.box_iou(bbox_tlist_anns, bbox_tlist_model)
```



- Essentially, it calculates the Intersection over Union (IoU) between all predicted bounding boxes of the ground truth and all predicted bounding boxes of the model. This process returns a tensor with the shape (number of ground truth bounding boxes, number of model's predicted bounding boxes). Refer to the images below.

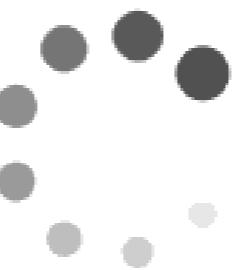


Here, we only need to obtain the maximum IoU for each ground truth bounding box and calculate the average value, so we use the following code

PYTHON

```
1 # After obtaining the maximum value for each predicted box of the annotation (see supplemental)
2 avg_iou = np.mean([t.cpu().detach().numpy().max() for t in iou]) # calculate the mean of IoU
```

! You might be curious whether using functions like `max()`, `mean()`, `sum()` will affect our results?



As we can see from the above image

- Using `sum()`, you may find that the value can exceed 1, which is not a reasonable range for IoU values.
- Using `max()`, it chooses, for each ground truth bounding box, the closest predicted bounding box from the model as its IoU. Then, we can obtain the maximum IoU values for `all predicted bounding boxes` of the ground truth and calculate the average to determine the overall IoU.
- Using `mean()` poses a problem as the IoU calculation will never be 1. This is because it considers the IoUs of other bounding boxes, which lowers the overall IoU. For instance, if the ground truth has two bounding boxes `[A1,A2]`, and the model also predicts two `[B1,B2]`, it's clear that B1 predicts A1, and B2 predicts A2, and the model predicts accurately. However, using mean will incorrectly consider B1 as A2 and B2 as A1, which is wrong, and these pairs have low IoUs. Thus, using `mean()` in this way will unjustly lower the IoU, making it unreasonable.

⌚ Author: [Shannon Hung](#)

↗ Link: <https://shannonhung.github.io/en/posts/coco-object-diagnose/>

❗ Copyright Notice: All articles in this blog are licensed under [CC BY-NC-SA 4.0](#) unless stating additionally.

Machine Learning



Related Articles

2023-11-16
Twitter Dataset - Using LSTM to predict the emotion of the article

2023-09-30
MAC OS - PyTorch on Mac OS with GPU support

2023-10-31
Flower102 Dataset - Using Transfer Learning to train + Usin...

2024-09-04
Windows - A Step-by-Step Guide to Enable CUDA GPU on...

2023-10-24
CIFAR10 Dataset - Using Pytorch to build CNN + activate GPU +...



Comment

0 comments

Anonymous ▾



Leave a comment

ⓘ Markdown is supported

Login with GitHub

Preview

Be the first person to leave a comment!

